


# Binary Search Trees


Chapter 11



# Binary Search




- Binary search can perform nearest neighbor queries on an ordered map that is implemented with an array, sorted by key
  - at each step, the number of candidate items is halved
  - terminates after  $O(\log n)$  steps
- Example: **find(7)**



# Binary Search Trees

- A binary search tree is a binary tree storing keys (or key-value entries) at its internal nodes and satisfying the following property:
  - Let  $u$ ,  $v$ , and  $w$  be three nodes such that  $u$  is in the left subtree of  $v$  and  $w$  is in the right subtree of  $v$ . We have  $\text{key}(u) \leq \text{key}(v) \leq \text{key}(w)$
- External nodes do not store items
- An inorder traversal of a binary search tree visits the keys in increasing order




# Search

- To search for a key  $k$ , we trace a downward path starting at the root
- The next node visited depends on the comparison of  $k$  with the key of the current node
- If we reach a leaf, the key is not found
- Example:  $\text{get}(4)$ :
  - Call  $\text{TreeSearch}(4, \text{root})$



**Algorithm  $\text{TreeSearch}(k, v)$**

```
if  $T.\text{isExternal}(v)$ 
    return  $v$ 
if  $k < \text{key}(v)$ 
    return  $\text{TreeSearch}(k, \text{left}(v))$ 
else if  $k = \text{key}(v)$ 
    return  $v$ 
else {  $k > \text{key}(v)$  }
    return  $\text{TreeSearch}(k, \text{right}(v))$ 
```





# Insertion

- To perform operation  $\text{put}(k, o)$ , we search for key  $k$  (using TreeSearch)
- Assume  $k$  is not already in the tree, and let  $w$  be the leaf reached by the search
- We insert  $k$  at node  $w$  and expand  $w$  into an internal node
- Example: insert 5





# Deletion

- To perform operation  $\text{remove}(k)$ , we search for key  $k$
- Assume key  $k$  is in the tree, and let  $v$  be the node storing  $k$
- If node  $v$  has a leaf child  $w$ , we remove  $v$  and  $w$  from the tree with operation  $\text{removeExternal}(w)$ , which removes  $w$  and its parent
- Example: remove 4





# Deletion (cont.)

- We consider the case where the key  $k$  to be removed is stored at a node  $v$  whose children are both internal
  - we find the internal node  $w$  that follows  $v$  in an inorder traversal
  - we copy  $\text{key}(w)$  into node  $v$
  - we remove node  $w$  and its left child  $z$  (which must be a leaf) by means of operation `removeExternal(z)`
- Example: remove 3




# Deletion (cont.)

- We consider the case where the key  $k$  to be removed is stored at a node  $v$  whose children are both internal
  - we find the internal node  $w$  that precedes  $v$  in an inorder traversal
  - we copy  $\text{key}(w)$  into node  $v$
  - we remove node  $w$  and its right child  $z$  (which must be a leaf) by means of operation `removeExternal(z)`
- Example: remove 8




# Performance

- Consider an ordered map with  $n$  items implemented by means of a binary search tree of height  $h$ 
  - the space used is  $O(n)$
  - methods `get`, `put` and `remove` take  $O(h)$  time
- The height  $h$  is  $O(n)$  in the worst case and  $O(\log n)$  in the best case
- On average, a binary tree generated from random insertions and removals of keys has expected height  $O(\log n)$




# AVL Trees



# AVL Tree Definition

- ❑ Inventors - Adel'son, Vel'skii and Landis
- ❑ AVL trees are balanced
- ❑ An AVL Tree is a **binary search tree** such that for every internal node  $v$  of  $T$ , the heights of the children of  $v$  can differ by at most 1
- ❑ Height-balance property



An example of an AVL tree where the heights are shown next to the nodes


# Height of an AVL Tree

Fact: The height of an AVL tree storing  $n$  keys is  $O(\log n)$ .

Proof (by induction): Let us bound  $m(h)$ : the minimum number of internal nodes of an AVL tree of height  $h$ .

- We easily see that  $m(1) = 1$  and  $m(2) = 2$
- For  $n > 1$ , an AVL tree of height  $h$  contains the root node, one AVL subtree of height  $h-1$  and another of height  $h-2$  (or  $h-1$ , but will we pick  $h-1$ ?).
- That is,  $m(h) = 1 + m(h-1) + m(h-2)$
- Knowing  $m(h-1) \geq m(h-2)$ , we get  $m(h) > 2m(h-2)$ . So  
 $m(h) > 2m(h-2)$ ,  $m(h) > 4m(h-4)$ ,  $m(h) > 8m(h-6)$ , ... (by induction),  
 $m(h) > 2^i m(h-2i)$
- Solving the base case we get:  $m(h) > 2^{h-1/2}$
- Taking logarithms:  $h < 2\log m(h)+1$
- Thus, the height of an AVL tree is  $O(\log n)$

Exercise: Stricter bound- AVL tree with  $n$  nodes has height at most  $\log_{1.63} n$



# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .

# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children.

# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children. (proof by contradiction)
    - ◆ all levels 1 to  $k-1$  are full (all nodes are present)
    - ◆ Hence the tree has at least  $2^{k-1}$  nodes

# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children. (proof by contradiction)
    - ◆ all levels 1 to  $k-1$  are full (all nodes are present)
    - ◆ Hence the tree has at least  $2^{k-1}$  nodes

# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
    - ◆ hence tree has at most  $2^{2k-1}$  nodes.
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children.
    - ◆ all levels 1 to  $k-1$  are full
    - ◆ hence tree has at least  $2^{k-1}$  nodes

# Structure of AVL Tree

- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
    - ◆ hence tree has at most  $2^{2k-1}$  nodes.
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children.
    - ◆ all levels 1 to  $k-1$  are full
    - ◆ hence tree has at least  $2^{k-1}$  nodes
  - Thus  $2^{k-1} \leq n \leq 2^{2k-1}$

# Structure of AVL Tree


- Consider an AVL tree on  $n$  nodes and the leaf that is closest to the root is at level  $k$ . Then
  - The height of the tree is at most  $2k-1$ .
    - ◆ hence tree has at most  $2^{2k-1}$  nodes.
  - All nodes at levels  $1, 2, \dots, k-2$  have 2 children.
    - ◆ all levels 1 to  $k-1$  are full
    - ◆ hence tree has at least  $2^{k-1}$  nodes
  - Thus  $2^{k-1} \leq n \leq 2^{2k-1}$
- Substituting  $h = 2k-1$ , we get  $2^{(h-1)/2} \leq n \leq 2^h$

# Summary - Structure of AVL Tree


- In an AVL tree of height  $h$ , the leaf closest to the root is at level  $(h+1)/2$
- On the first  $(h-1)/2$  levels, the AVL tree is a complete binary tree
  - thins out after  $(h-1)/2$  level
- $2^{(h-1)/2} \leq$  number of nodes  $\leq 2^h$

# Insertion

- ❑ Insertion is as in a binary search tree
- ❑ Always done by expanding an external node.
- ❑ Example: insert 54




before insertion




after insertion

# Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.




before insertion




after insertion

# Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.




before insertion




after insertion

# Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.
- The only nodes whose heights can increase are the ancestors of inserted node.




before insertion



after insertion


# Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.
- The only nodes whose heights can increase are the ancestors of inserted node.
- If the insertion causes T to become unbalanced, then some ancestor of the inserted node would have a height imbalance.




# Insertion

- Inserting a node into an AVL tree changes the height of some of the nodes in T.
- The only nodes whose heights can increase are the ancestors of inserted node.
- If the insertion causes T to become unbalanced, then some ancestor of the inserted node would have a height imbalance.
- We travel up the tree from the inserted node (v), until we find the first node (x) such that its grandparent (z) is unbalanced.
- Let y be the parent of node x.




# Insertion

- To rebalance the subtree rooted at z, we must perform a rotation.




AVL Trees




# Rotations

- Rotation is a way of locally reorganizing a BST.
- Let  $x, y$  be two nodes such that  $y=\text{parent}(x)$ .
- $\text{Keys}(T_1) < \text{Key}(x) < \text{Keys}(T_2) < \text{Key}(y) < \text{Keys}(T_3)$




# Rotations

- Rotation is a way of locally reorganizing a BST.
- Let  $x, y$  be two nodes such that  $y=\text{parent}(x)$ .
- $\text{Keys}(T_1) < \text{Key}(x) < \text{Keys}(T_2) < \text{Key}(y) < \text{Keys}(T_3)$



# Restructuring – Single Rotation


- Rotations to make y the top-most node



- Insertion happens in subtree  $T_2$
- $H(T_2)$  increases from  $h$  to  $h+1$
- $x$  remains balanced
  - $H(T_3) = h$  or  $h+1$  or  $h+2$ 
    - $h+2$  – then  $x$  is originally unbalanced
    - $h+1$  – then height of  $x$  does not increase –  $z$  is balanced
    - therefore  $H(T_3) = h$
- $H(x)$  increases from  $h+1$  to  $h+2$


# Restructuring – Single Rotation

- Rotations to make y the top-most node
- Insertion happens in subtree  $T_2$
- y remains balanced
  - $H(T_1) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆  $h+3$  – y is originally unbalanced
    - ◆  $h+2$  – height of y does not increase – z is balanced
    - ◆ therefore  $H(T_1) = h+1$
  - $H(y) = h+2$  to  $h+3$ .




# Restructuring – Single Rotation

- Rotations to make y the top-most node
- Insertion happens in subtree  $T_2$
- y remains balanced
  - $H(T_1) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆  $h+3$  – y is originally unbalanced
    - ◆  $h+2$  – height of y does not increase – z is balanced
    - ◆ therefore  $H(T_1) = h+1$
- $H(y) = h+2$  to  $h+3$ .
- z is unbalanced
  - $H(T_0) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆ Since originally z was balanced
    - ◆  $H(T_0) = h+1$
- $H(z) = h+3$  to  $h+4$





# Restructuring – Single Rotation

- Rotations to make y the top-most node
- Insertion happens in subtree  $T_2$
- y remains balanced
  - $H(T_1) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆  $h+3$  – y is originally unbalanced
    - ◆  $h+2$  – height of y does not increase – z is balanced
    - ◆ therefore  $H(T_1) = h+1$
  - $H(y) = h+2$  to  $h+3$ .
  - z is unbalanced
    - $H(T_0) = h+1$  or  $h+2$  or  $h+3$ 
      - ◆ Since originally z was balanced
      - ◆  $H(T_0) = h+1$
    - $H(z) = h+3$  to  $h+4$




# Restructuring – Single Rotation

- Rotations to make  $y$  the top-most node




# Restructuring – Single Rotation

- Rotations to make y the top-most node




# Restructuring – Single Rotation

- Rotations to make  $y$  the top-most node




# Restructuring – Single Rotation

- Rotations to make y the top-most node




- Symmetric to the previous scenario




# Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Insertion happens in subtree  $T_1$ 
  - $H(T_1) = h$  to  $h+1$
  - x is balanced
    - $H(T_2) = h$  or  $h+1$  or  $h+2$ 
      - ◆  $h+2$  – x is originally unbalanced
      - ◆  $h+1$  – no increase in height of x – z remains balanced
      - ◆ therefore  $H(T_2) = h$
  - $H(x) = h+1$  to  $h+2$



# Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- Insertion happens in subtree  $T_1$
- y remains balanced
  - $H(T_3) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆  $h+3$  – y is originally unbalanced
    - ◆  $h+2$  – no increase in height of y – z remains balanced
    - ◆ therefore  $H(T_3) = h+1$




- $H(y) = h+2$  to  $h+3$

# Restructuring – Double Rotation

- Perform rotations around  $y$  and  $z$  to make  $x$  the top-most node.
- Insertion happens in subtree  $T_1$
- $y$  remains balanced

- $H(T_3) = h+1$  or  $h+2$  or  $h+3$ 
  - ◆  $h+3$  –  $y$  is originally unbalanced
  - ◆  $h+2$  – no increase in height of  $y - z$  remains balanced
  - ◆ therefore  $H(T_3) = h+1$

- $H(y) = h+2$  to  $h+3$
- $z$  is unbalanced
  - $H(T_0) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆ since  $z$  was originally balanced
    - ◆  $H(T_0) = h+1$
- $H(z) = h+3$  to  $h+4$




# Restructuring – Double Rotation

- Perform rotations around  $y$  and  $z$  to make  $x$  the top-most node.
- Insertion happens in subtree  $T_1$
- $y$  remains balanced

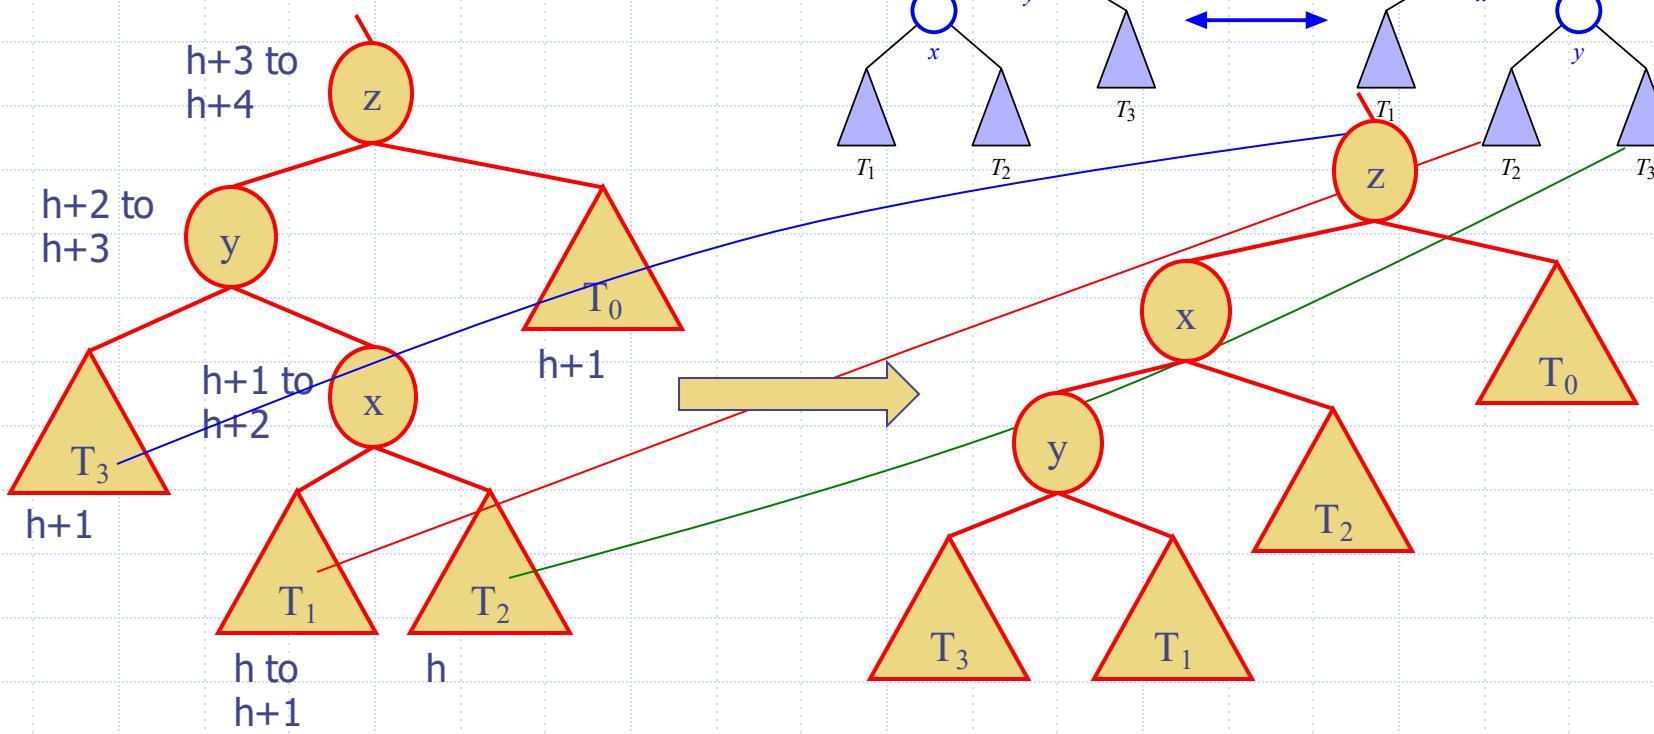
- $H(T_3) = h+1$  or  $h+2$  or  $h+3$ 
  - ◆  $h+3$  –  $y$  is originally unbalanced
  - ◆  $h+2$  – no increase in height of  $y - z$  remains balanced
  - ◆ therefore  $H(T_3) = h+1$

- $H(y) = h+2$  to  $h+3$
- $z$  is unbalanced
  - $H(T_0) = h+1$  or  $h+2$  or  $h+3$ 
    - ◆ since  $z$  was originally balanced
    - ◆  $H(T_0) = h+1$
- $H(z) = h+3$  to  $h+4$




# Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- First step
  - rotate x and y



# Restructuring – Double Rotation


- Perform rotations around  $y$  and  $z$  to make  $x$  the top-most node.
- First step
  - rotate  $x$  and  $y$



# Restructuring – Double Rotation


- Perform rotations around y and z to make x the top-most node.

- Second step
  - rotate x and z



# Restructuring – Double Rotation


- Perform rotations around  $y$  and  $z$  to make  $x$  the top-most node.
- Second step
  - rotate  $x$  and  $z$



AVL Trees


# Restructuring – Double Rotation

- Perform rotations around  $(x,y)$  and  $(x,z)$  to make  $x$  the top-most node.





# Restructuring – Double Rotation

- Perform rotations around y and z to make x the top-most node.
- symmetric to the previous configuration




# Insertion Example, continued




# Removal

- Removal begins as in a binary search tree, which means the node removed will become an empty external node. Its parent, w, may cause an imbalance.
- Example:




# Rebalancing after a Removal

- Let  $z$  be the first unbalanced node encountered while travelling up the tree from  $w$ . Also, let  $y$  be the child of  $z$  with the larger height, and let  $x$  be the child of  $y$  with the larger height
- We perform a **trinode restructuring** to restore balance at  $z$
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of  $T$  is reached




# Deletion




- Suppose deletion happens in subtree  $T_0$  and its height reduces from  $h$  to  $h-1$
- $z$  was originally balanced and now unbalanced
  - $H(y) = h+1$
  - $H(z) = h+2$
- $x$  has larger height than  $T_1$ 
  - $H(x) = h$
- $y$  is balanced
  - $H(T_1) = h$  or  $h-1$

# Deletion




- Suppose deletion happens in subtree  $T_0$  and its height reduces from  $h$  to  $h-1$
- $x$  is balanced
  - $H(T_2), H(T_3)$  is  $h-1$  or  $h-2$
  - However both  $T_2$  and  $T_3$  cannot have height  $h-2$

# Deletion – Restructuring – Single Rotation




# Deletion – Restructuring – Single Rotation

- After rotation the height of the subtree might be less than 1




# Deletion – Restructuring – Double Rotation



- Suppose deletion happens in subtree  $T_0$  and its height reduces from  $h$  to  $h-1$
- $z$  was originally balanced and now unbalanced
  - $H(y) = h+1$
  - $H(z) = h+2$
- $x$  has larger height than  $T_3$ 
  - $H(x) = h$
- $y$  is balanced
  - $H(T_3) = h-1$


# Deletion – Restructuring – Double Rotation




- Suppose deletion happens in subtree  $T_0$  and its height reduces from  $h$  to  $h-1$
- $x$  remains balanced
  - $H(T_2), H(T_3)$  is  $h-1$  or  $h-2$
  - However both  $T_2$  and  $T_3$  cannot have height  $h-2$

# Deletion – Restructuring – Double Rotation

## Step 1




## Rotate $(x,y)$





# Deletion – Restructuring – Double Rotation

## □ Step 1




## □ Rotate $(x,y)$




AVL Trees

# Deletion – Restructuring – Double Rotation

## □ Step 2




## □ Rotate $(x,z)$



# Deletion – Restructuring – Double Rotation


- Step 2



- Rotate  $(x,z)$

- After rotation the height of the subtree is less than 1

# Rebalancing after a Removal



# AVL Tree Performance

- AVL tree storing  $n$  items

- The data structure uses  $O(n)$  space
- A single restructuring takes  $O(1)$  time
  - ◆ using a linked-structure binary tree
- Searching takes  $O(\log n)$  time
  - ◆ height of tree is  $O(\log n)$ , no restructures needed
- Insertion takes  $O(\log n)$  time
  - ◆ initial find is  $O(\log n)$
  - ◆ restructuring up the tree, maintaining heights is  $O(\log n)$
- Removal takes  $O(\log n)$  time
  - ◆ initial find is  $O(\log n)$
  - ◆ restructuring up the tree, maintaining heights is  $O(\log n)$

