

Bad Character Rule (BCR) (3)

- On a mismatch, shift the pattern to the right until the first occurrence of the mismatched character in P.
- Still $O(n m)$ worst case running time:

T: aaaaaaaaaaaaaaaaaaaaaaaaaaaaa

P: abaaaa

Good Suffix Rule (GSR) (1)

- We want to use the knowledge of the matched characters in the pattern's suffix.
- If we matched S characters in T , what is (if exists) the smallest shift in P that will align a sub-string of P of the same S characters ?

Good Suffix Rule (GSR) (2)

- Example 1 – how much to move:

T: bbacdcbabcddcdaddaaabc**cb**

P: cabbabdbab

cabbabdbab

Good Suffix Rule (GSR) (3)

- Example 2 – what if there is no alignment:



T: bbacdcbaabcbbabdbabcaabcccb

P: bcbabdbabc

bcbabdbabc

Good Suffix Rule (GSR) (4)

- We mark the matched sub-string in T with t and the mismatched char with x
 1. In case of a mismatch: shift right until the first occurrence of t in P such that the next char y in P holds $y \neq x$
 2. Otherwise, shift right to the largest prefix of P that aligns with a suffix of t .

Good Suffix Rule (GSR) (5)

- $L(i)$ – The biggest index j , such that $j < n$ and prefix $P[1..j]$ contains suffix $P[i..n]$ as a suffix but not suffix $P[i-1..n]$

| i | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|--------|---|---|---|---|---|---|---|---|---|
| P | G | T | A | G | C | G | G | C | G |
| $L(i)$ | 0 | 0 | 0 | 0 | 0 | 0 | 6 | 0 | 7 |

Good Suffix Rule (GSR) (6)

- $l(i)$ – The length of the longest suffix of $P[i..n]$ that is also a prefix of P

| i | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|--------|---|---|---|---|---|---|---|---|---|
| P | G | T | A | G | C | G | G | C | G |
| $l(i)$ | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |

Good Suffix Rule (GSR) (7)

□ Putting it together

- If mismatch occurs at position n , shift P by 1
- If a mismatch occurs at position $i-1$ in P :
 - ◆ If $L(i) > 0$, shift P by $n - L(i)$
 - ◆ else shift P by $n - l(i)$
- If P was found, shift P by $n - l(2)$

Boyer Moore Algorithm

- Preprocess(P)
- $k := n$
- while ($k \leq m$) do
 - Match P and T from right to left starting at k
 - If a mismatch occurs: shift P right (advance k) by max(good suffix rule, bad char rule).
 - else, print the occurrence and shift P right (advance k) by the good suffix rule.

Boyer-Moore Algorithm Demo

GT_TA_TA_GC_TG_AT_CG_CG_GC_GT_AG_CG_GC_GA_A

GT_AG_TC_GG_GC_G

Bad Character Rule – shift by 7

Good Suffix Rule – shift by 0

Boyer-Moore Algorithm Demo

GTTATAGCTGATCGCGGGCGTAGCGGCGAA

GTAGCGGGCG t

Bad Character Rule – shift by 1

Good Suffix Rule – shift by 3 (9-6)

Pattern P – GTAGCGGGCG t – GCG

Prefixes of P

G

GT

GTA

GTAG

GTAGC

GTAGCG

Find the longest prefix of P which has t as the suffix.

Boyer-Moore Algorithm Demo

GT_TA_TA_GC_TG_AT_CGC_GGG_GCG_TAGC_GGG_GCG_AA
GT_AAGC_GGG_GCG_T t

Pattern P – GTAGCGGCG

Prefixes of P

G
GT
GTA
GTAG
GTAGC
GTAGCG

t – GCGGCG

GTAGCGG
GTAGCGGC
GTAGCGGGCG

Find the longest prefix of P which has t as the suffix.

Boyer-Moore Algorithm Demo

GTTATAGCTGAT**C**GCGGCGTAGCGGCGAA

GTAGCGGCG *t*

Bad Character Rule – shift by 3

Good Suffix Rule – shift by 8 (9-1)

Pattern P – GTAGCGGCG *t* – GCGGCG

suffixes of *t* **Prefixes of P**

G G

CG GT

GCG GTA

GCGG GTAG

GCGGC GTAGC

GCGGCG GTAGCG

Find the longest prefix of
P which is a suffix of *t*.

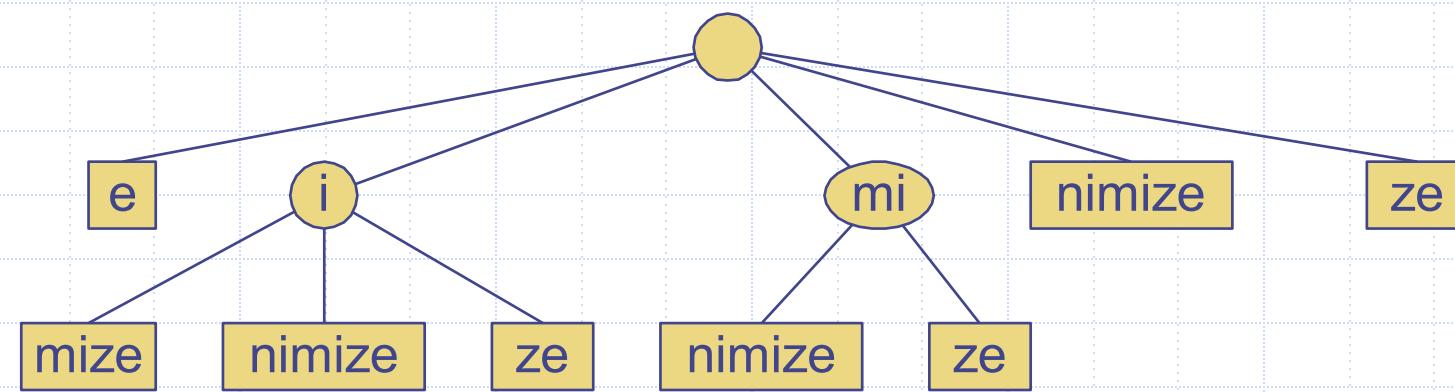
Boyer-Moore Algorithm Demo

GT~~TATAGCTGATCGC~~GGCGTAGCGGCGAA
GTAGCGGCCG

Boyer-Moore Algorithm Analysis

- Worst case
 - $O(nm)$ if the pattern does occur in the text
 - For patterns of small size, the algorithm might not be efficient.
- $O(n+m)$ if the pattern does not occur in the text

Tries

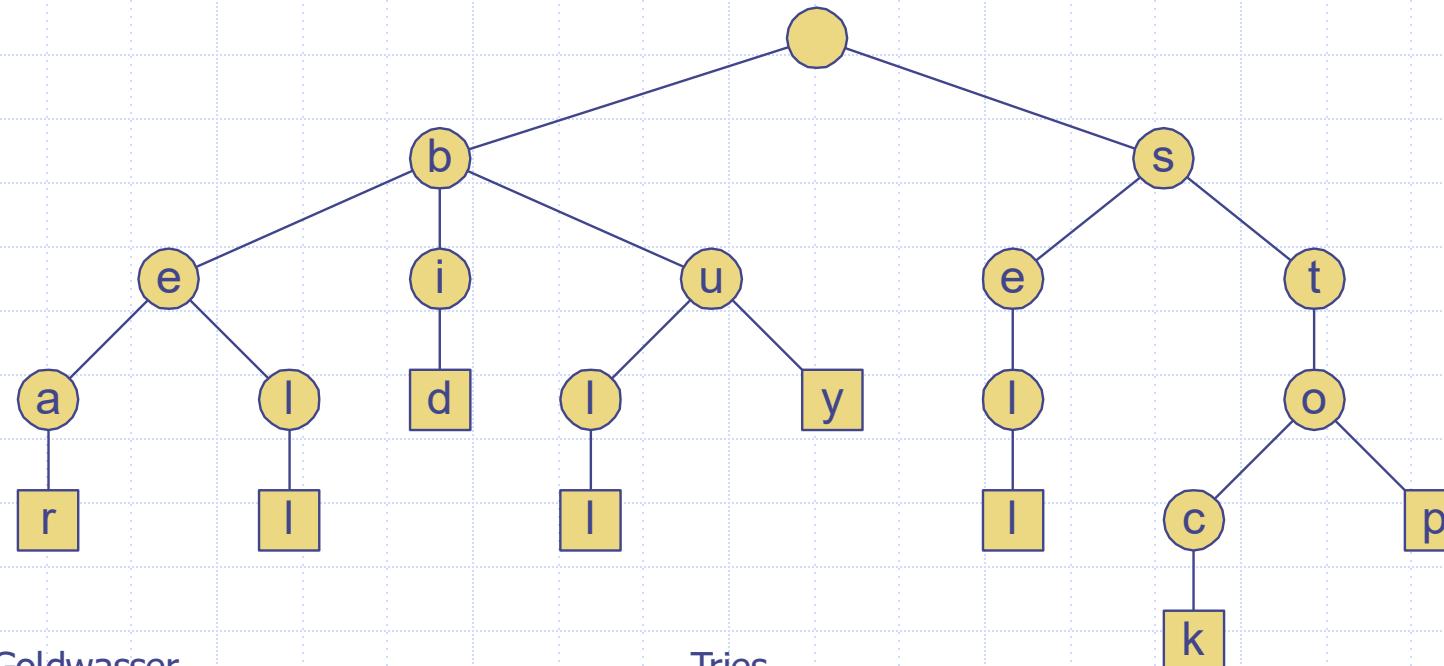


Preprocessing Strings

- Preprocessing the pattern speeds up pattern matching queries
- If the text is large, immutable and searched for often (e.g., works by Shakespeare), we may want to preprocess the text instead of the pattern
- A trie is a compact data structure for representing a set of strings, such as all the words in a text
 - A trie supports pattern matching queries in time proportional to the pattern size

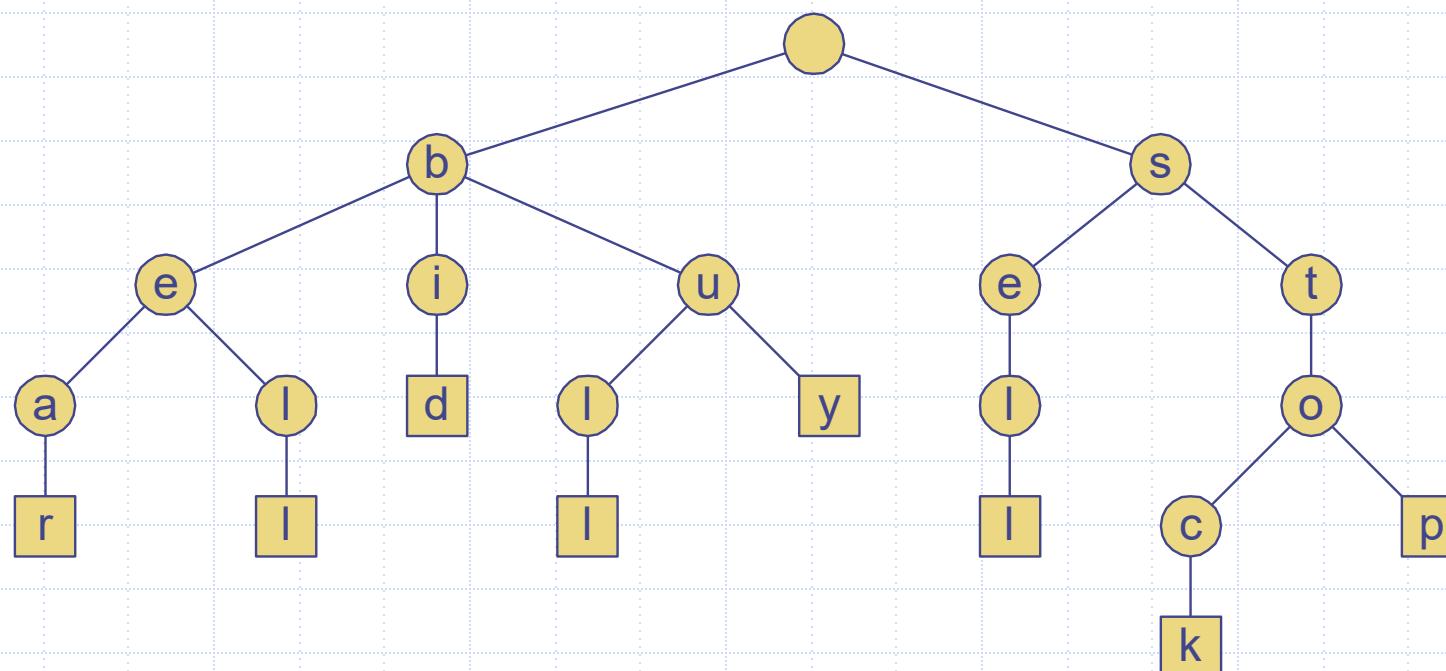
Standard Tries

- The standard trie for a set of strings S is an ordered tree such that:
 - Each node but the root is labeled with a character
 - The children of a node are alphabetically ordered
 - The paths from the external nodes to the root yield the strings of S
- Example: standard trie for the set of strings
 $S = \{ \text{bear}, \text{bell}, \text{bid}, \text{bull}, \text{buy}, \text{sell}, \text{stock}, \text{stop} \}$



Analysis of Standard Tries

- A standard trie uses $O(n)$ space and supports searches, insertions and deletions in time $O(dm)$, where:
 - n total size of the strings in S
 - m size of the string parameter of the operation
 - d size of the alphabet



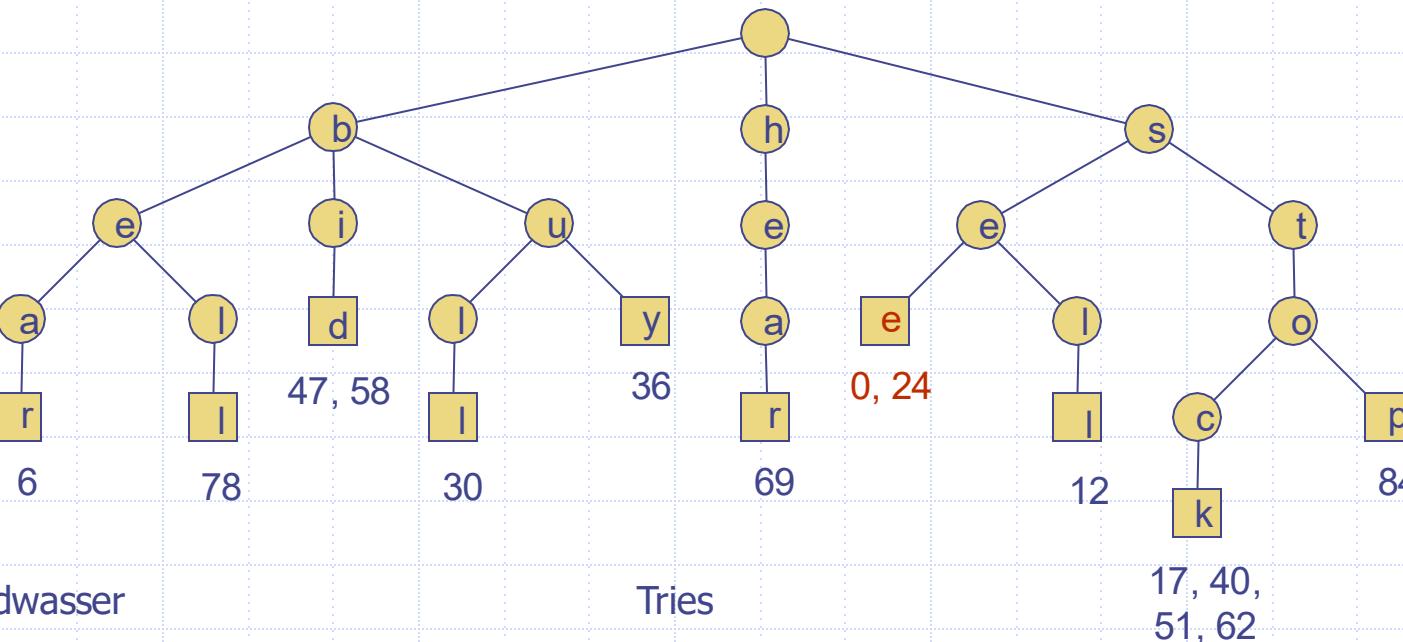
Application of Tries

- A standard trie supports the following operations on a processed text in $O(m)$ time, where m is the length of the string.
 - word matching: find the first occurrence of word X in the text.
 - prefix matching: find the first occurrence of the longest prefix of word X in the text.

Word Matching with a Trie

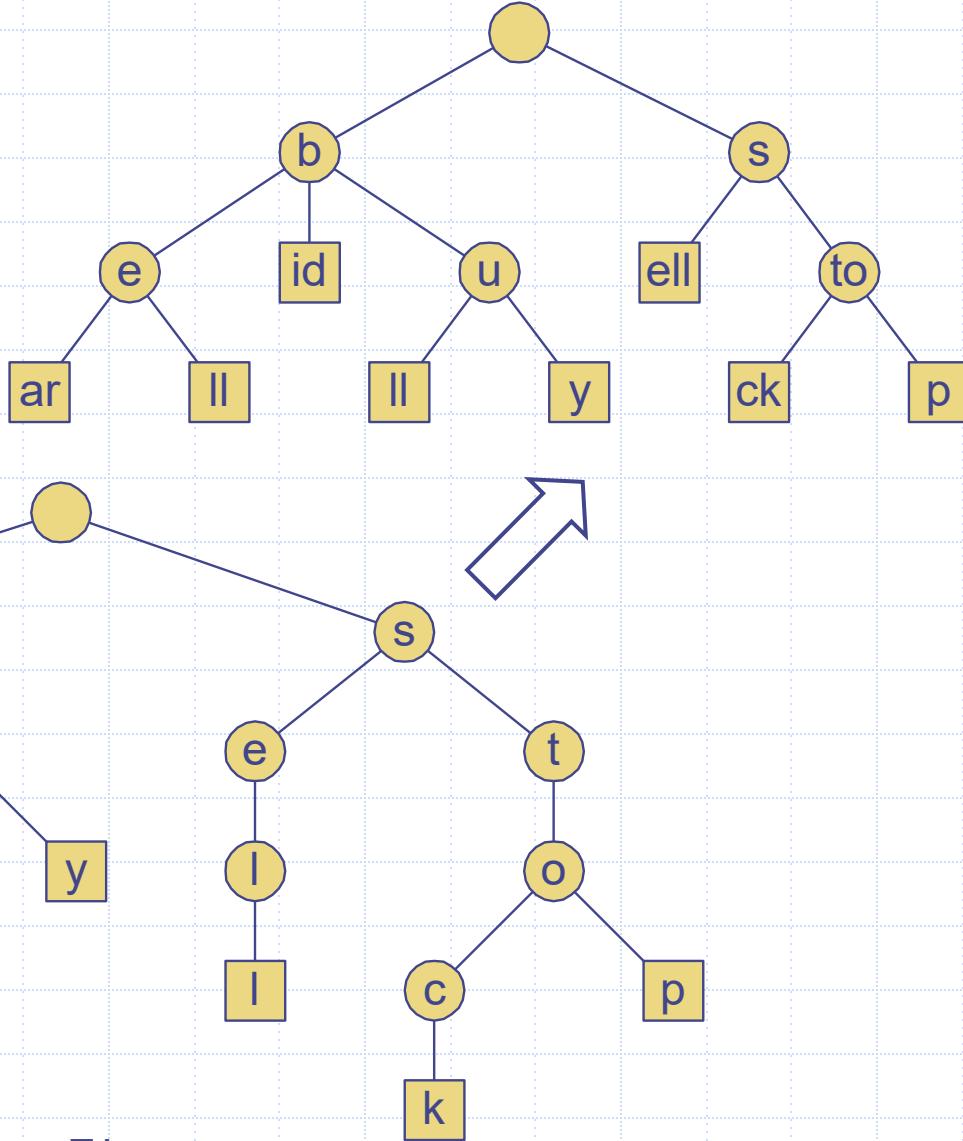
- ◆ insert the words of the text into trie
- ◆ Each leaf is associated w/ one particular word
- ◆ leaf stores indices where associated word begins ("see" starts at index 0 & 24, leaf for "see" stores those indices)

| | | | | | | | | | | | | | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| s | e | e | | a | b | e | a | r | ? | | s | e | l | l | s | t | o | c | k | ! | | | |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 |
| s | e | e | | a | b | u | l | l | ? | | b | u | y | | s | t | o | c | k | ! | | | |
| 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 | 32 | 33 | 34 | 35 | 36 | 37 | 38 | 39 | 40 | 41 | 42 | 43 | 44 | 45 | 46 | |
| b | i | d | | s | t | o | c | k | ! | | b | i | d | | s | t | o | c | k | ! | | | |
| 47 | 48 | 49 | 50 | 51 | 52 | 53 | 54 | 55 | 56 | 57 | 58 | 59 | 60 | 61 | 62 | 63 | 64 | 65 | 66 | 67 | 68 | | |
| h | e | a | r | | t | h | e | b | e | l | l | ? | | s | t | o | p | ! | | | | | |
| 69 | 70 | 71 | 72 | 73 | 74 | 75 | 76 | 77 | 78 | 79 | 80 | 81 | 82 | 83 | 84 | 85 | 86 | 87 | 88 | | | | |



Compressed Tries

- ❑ A compressed trie has internal nodes of degree at least two
- ❑ It is obtained from standard trie by compressing chains of "redundant" nodes
- ❑ ex. the "i" and "d" in "bid" are "redundant" because they signify the same word

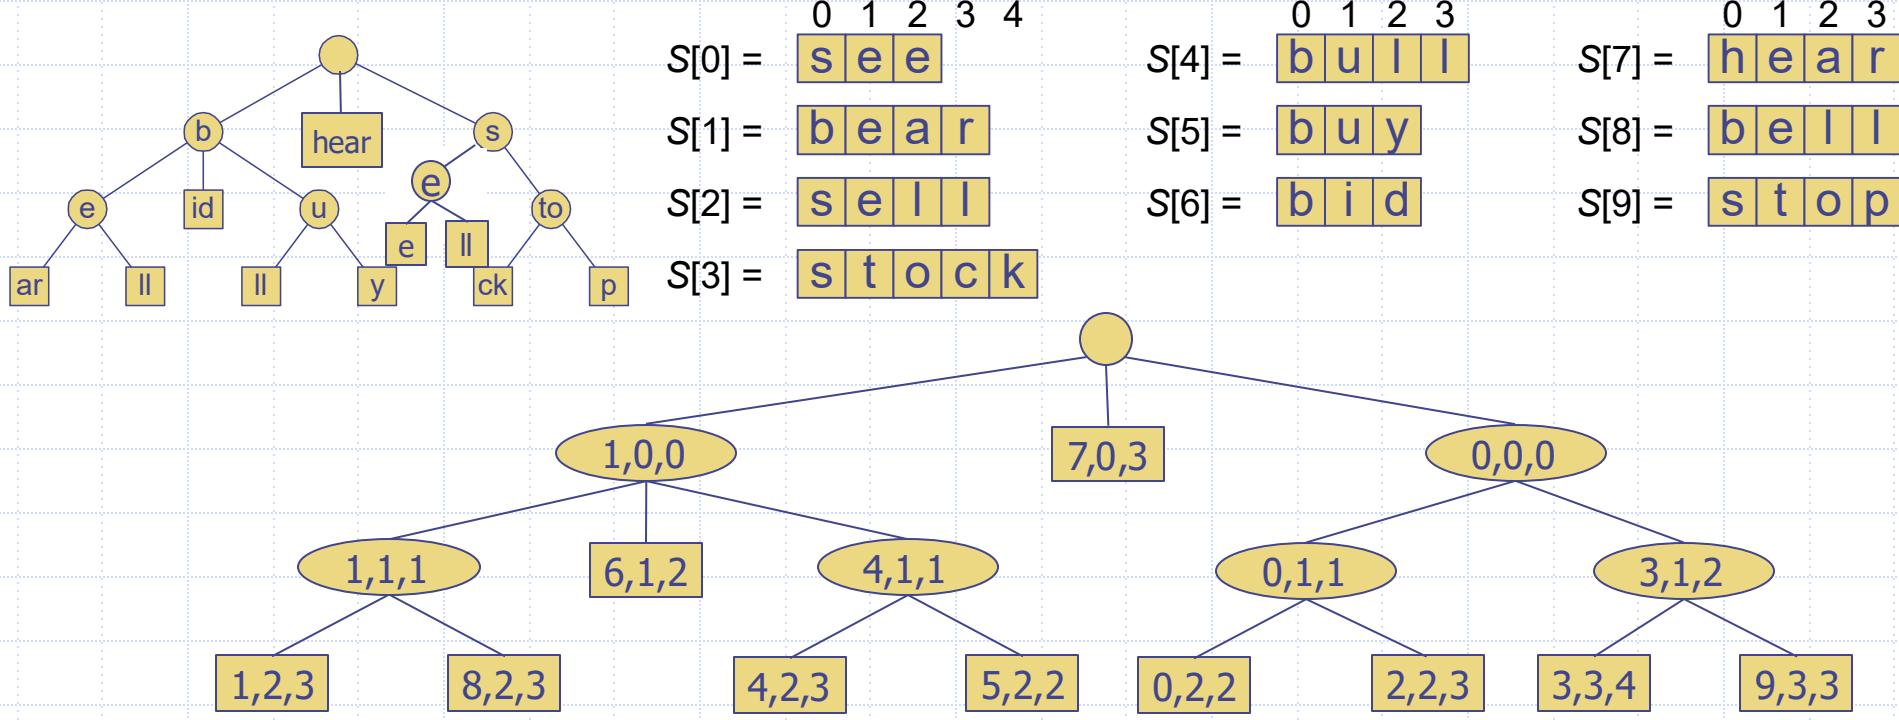


Why Compressed Tries?

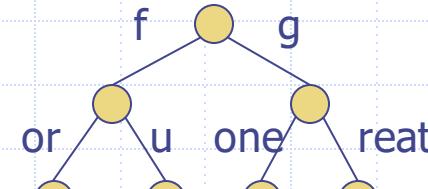
- A compressed trie storing a collection S of s strings from an alphabet of size d has the following properties
 - Every internal node of T has at least two children and at most d children
 - T has s leaf nodes
 - The number of internal nodes of T is $O(s)$

Compact Representation

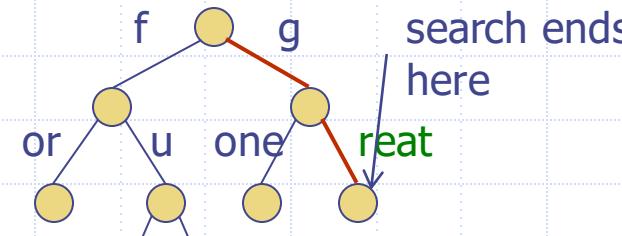
- Compact representation of a compressed trie for an array of strings:
 - Stores at the nodes ranges of indices instead of substrings
 - Uses $O(s)$ space, where s is the number of strings in the array
 - Serves as an auxiliary index structure



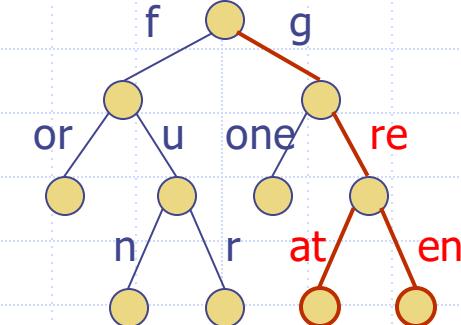
Insertion/Deletion in Compressed Trie



insert green

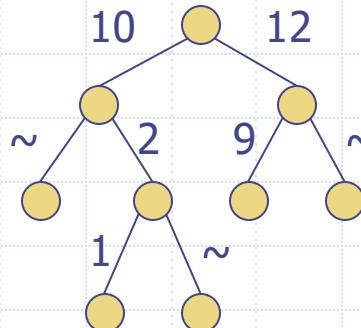


insert green



Application - Routing through Tries

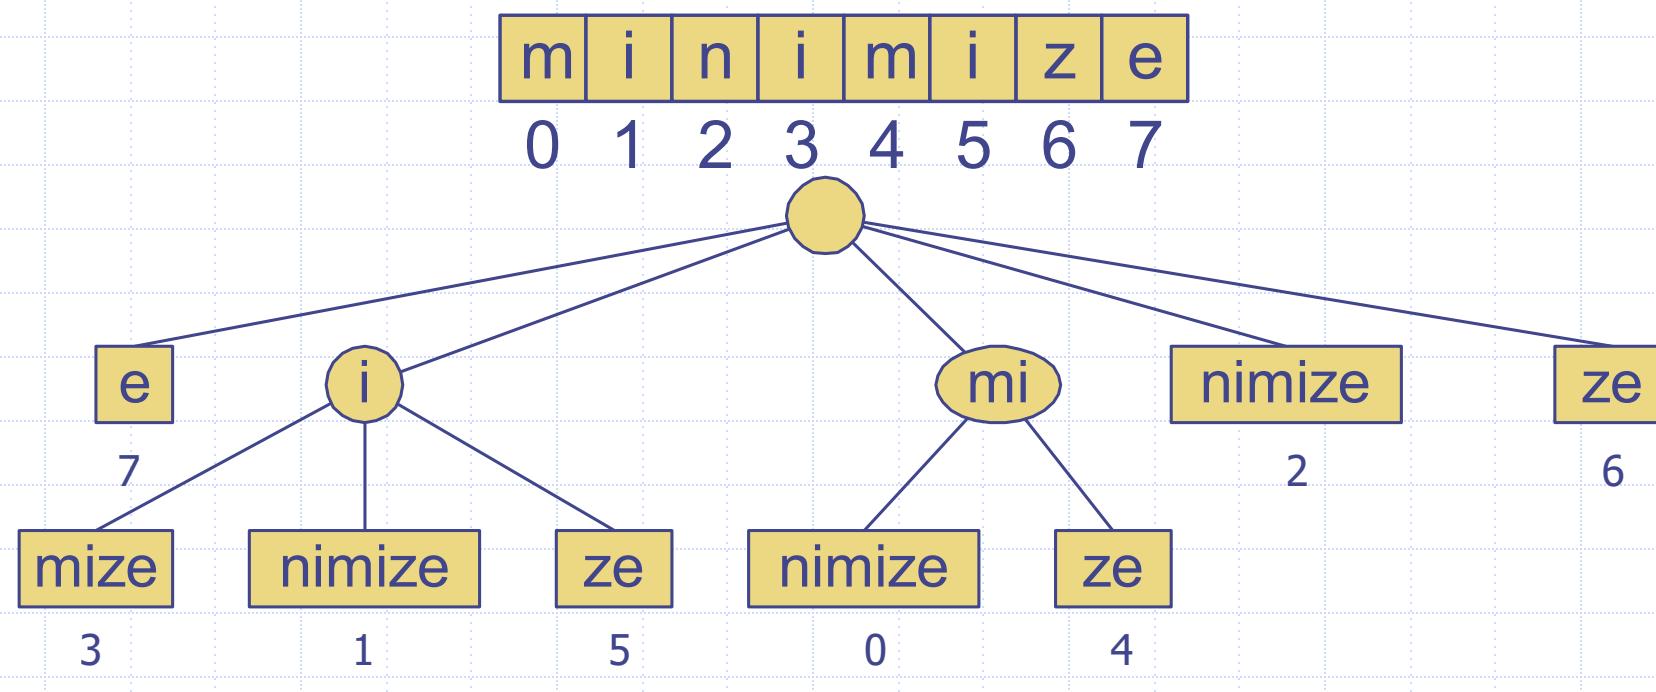
- Internet Routers maintain a Trie (table)
- It is not a lookup
 - forwards packets to its neighbors using IP prefix matching rules



Tries

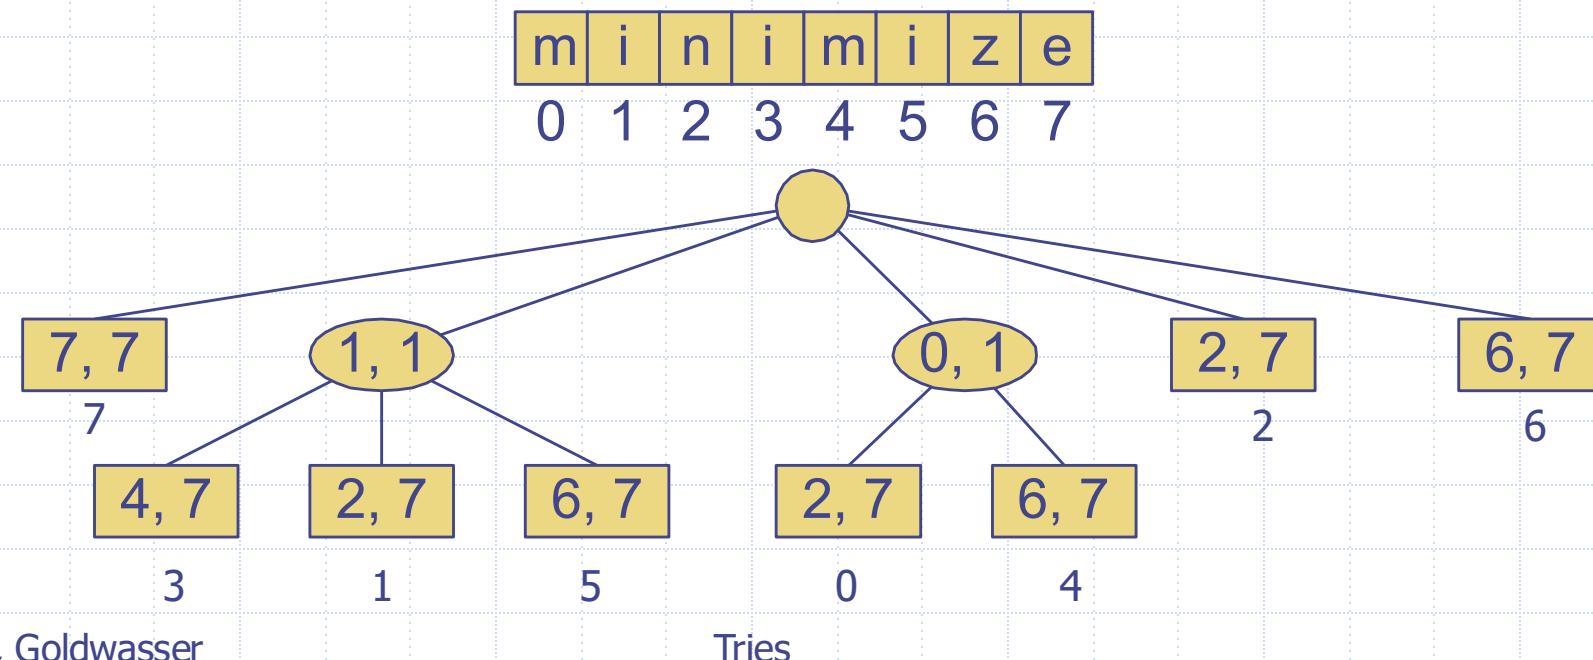
Suffix Trie

- The suffix trie of a string X is the compressed trie of all the suffixes of X
- Each leaf corresponds to a suffix of X



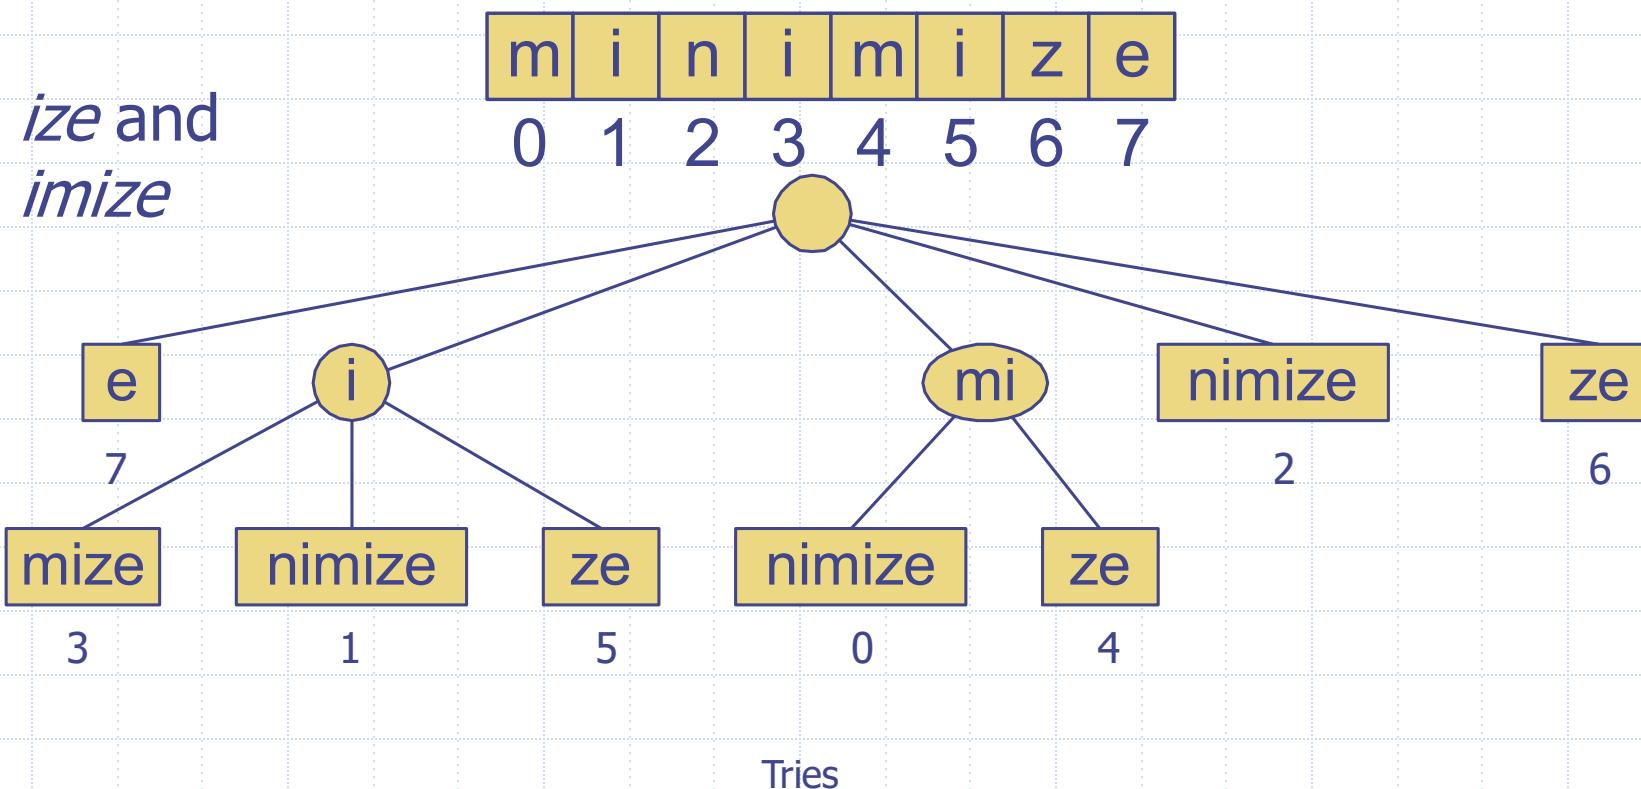
Analysis of Suffix Tries

- Compact representation of the suffix trie for a string X of size n from an alphabet of size d
 - Uses $O(n)$ space
 - Supports arbitrary pattern matching queries in X in $O(dm)$ time, where m is the size of the pattern
 - Can be constructed in $O(n)$ time



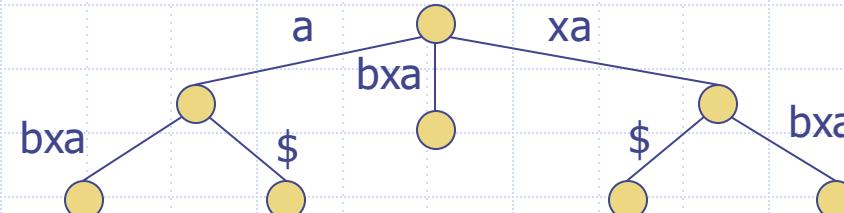
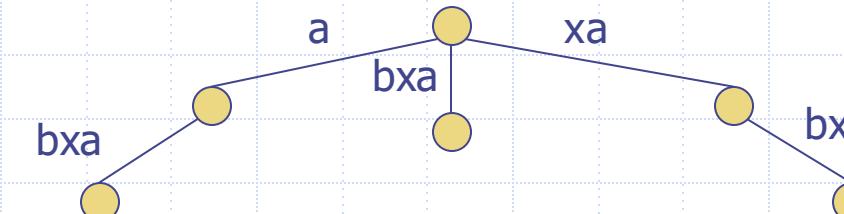
Prefix Matching using Suffix Trie

- If two suffixes have a same prefix, then their corresponding paths are the same at their beginning, and the concatenation of the edge labels of the mutual part is the prefix.



Suffix Trie

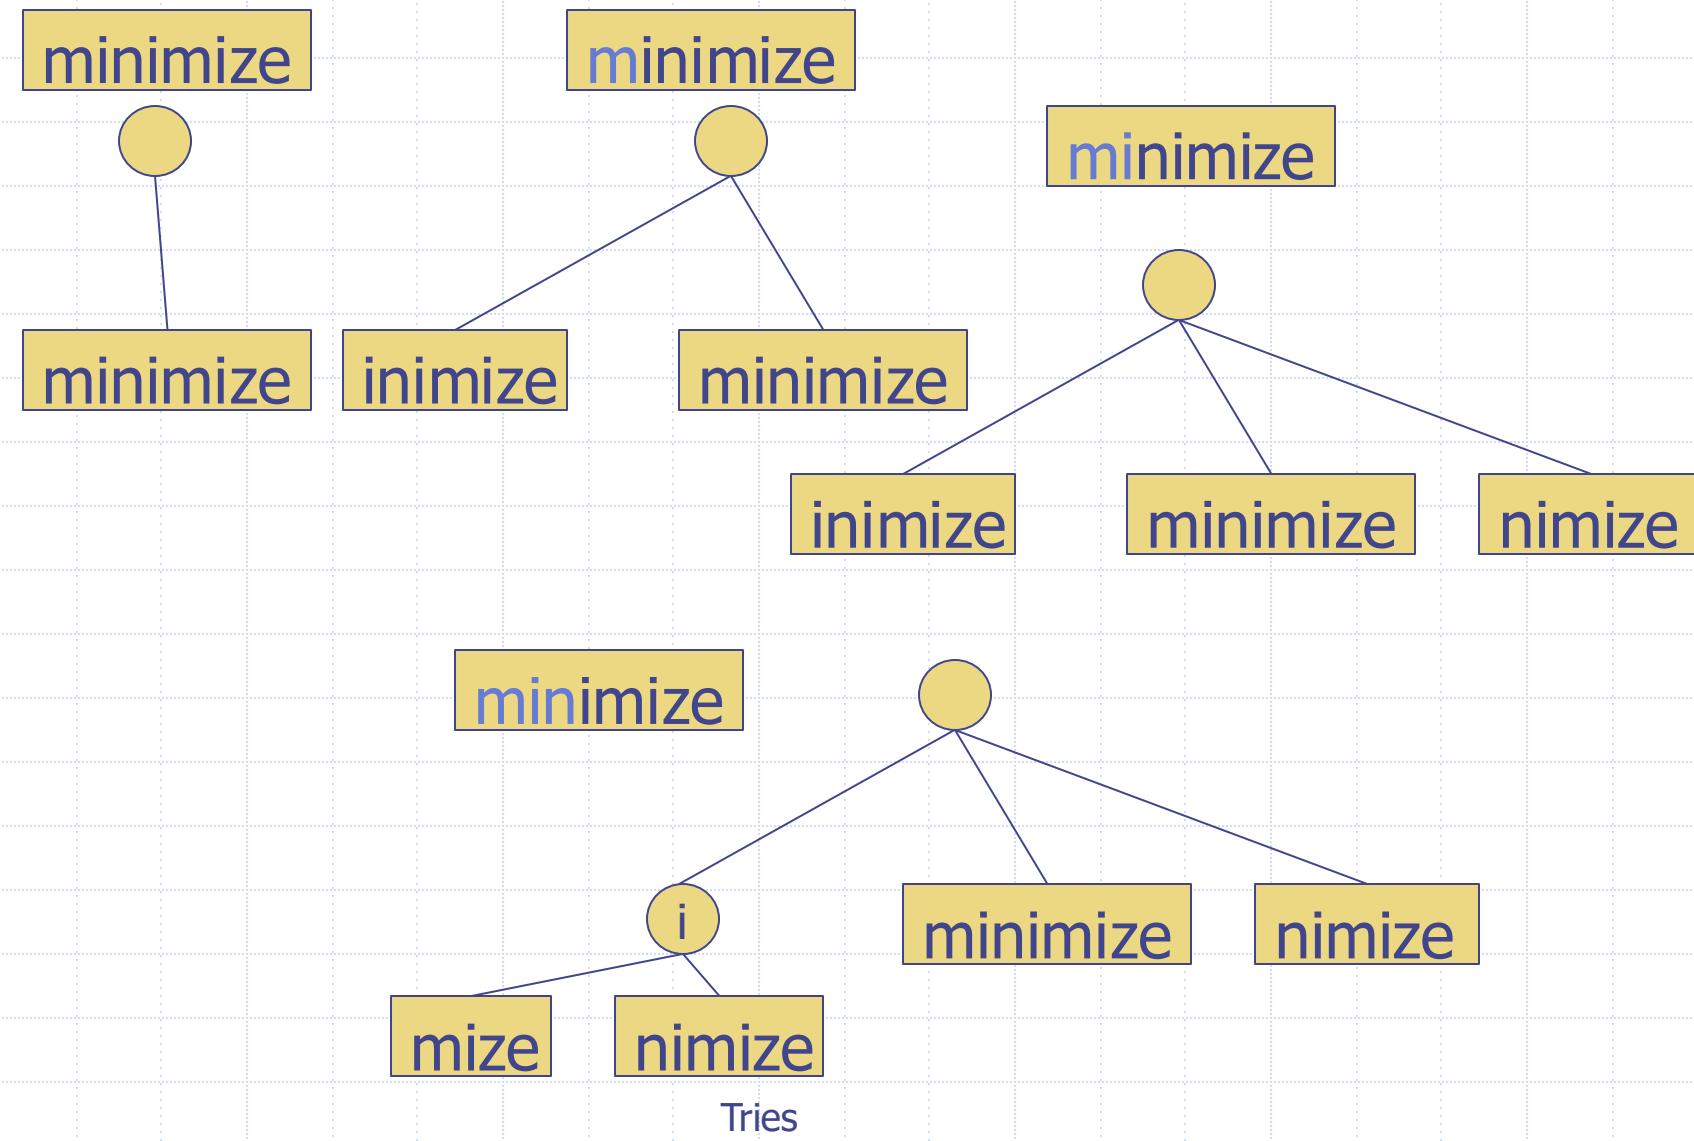
- Not all strings are guaranteed to have corresponding suffix trie.
- For example: xabxa



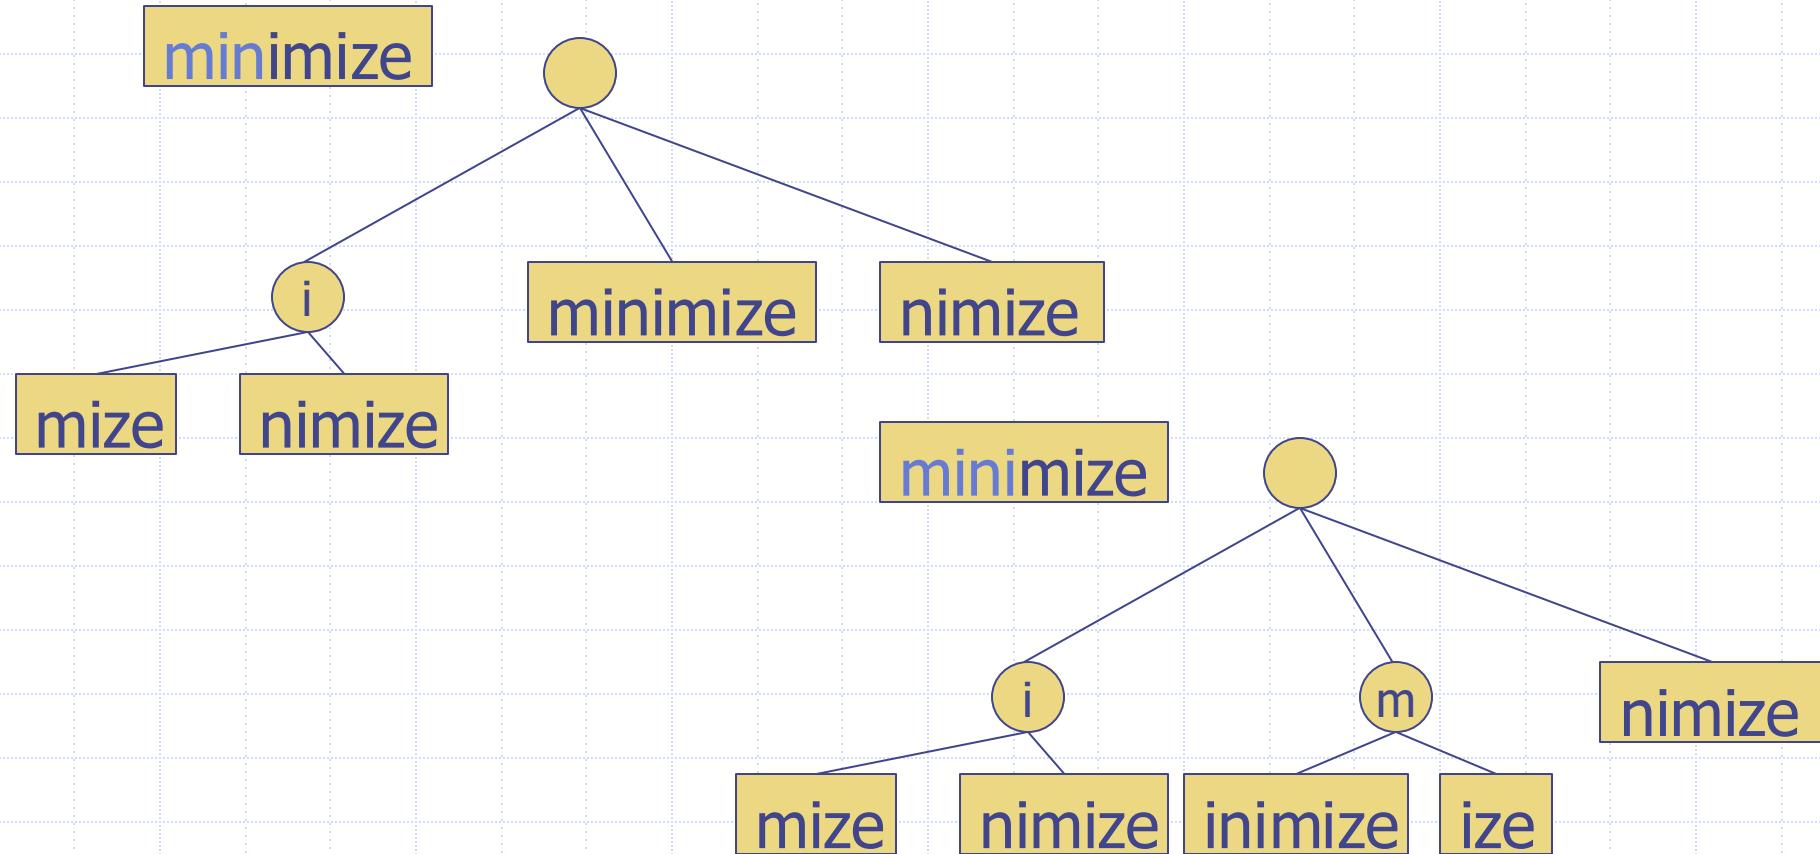
Constructing a Suffix Trie

- $S[1..n]$ is the string
- start with a single edge for S
- enter the edges for the suffix $S[i..n]$ where i goes from 2 to n
 - Starting at the root node find the longest part from the root whose label matches a prefix of $S[i..n]$. At some point, no further matches are possible
 - ◆ If the point is at a node, then denote this node by w
 - ◆ If it is in the middle of an edge, then insert a new node called w , at this point
 - ◆ create a new edge running from w to a new leaf labeled $S[i..n]$

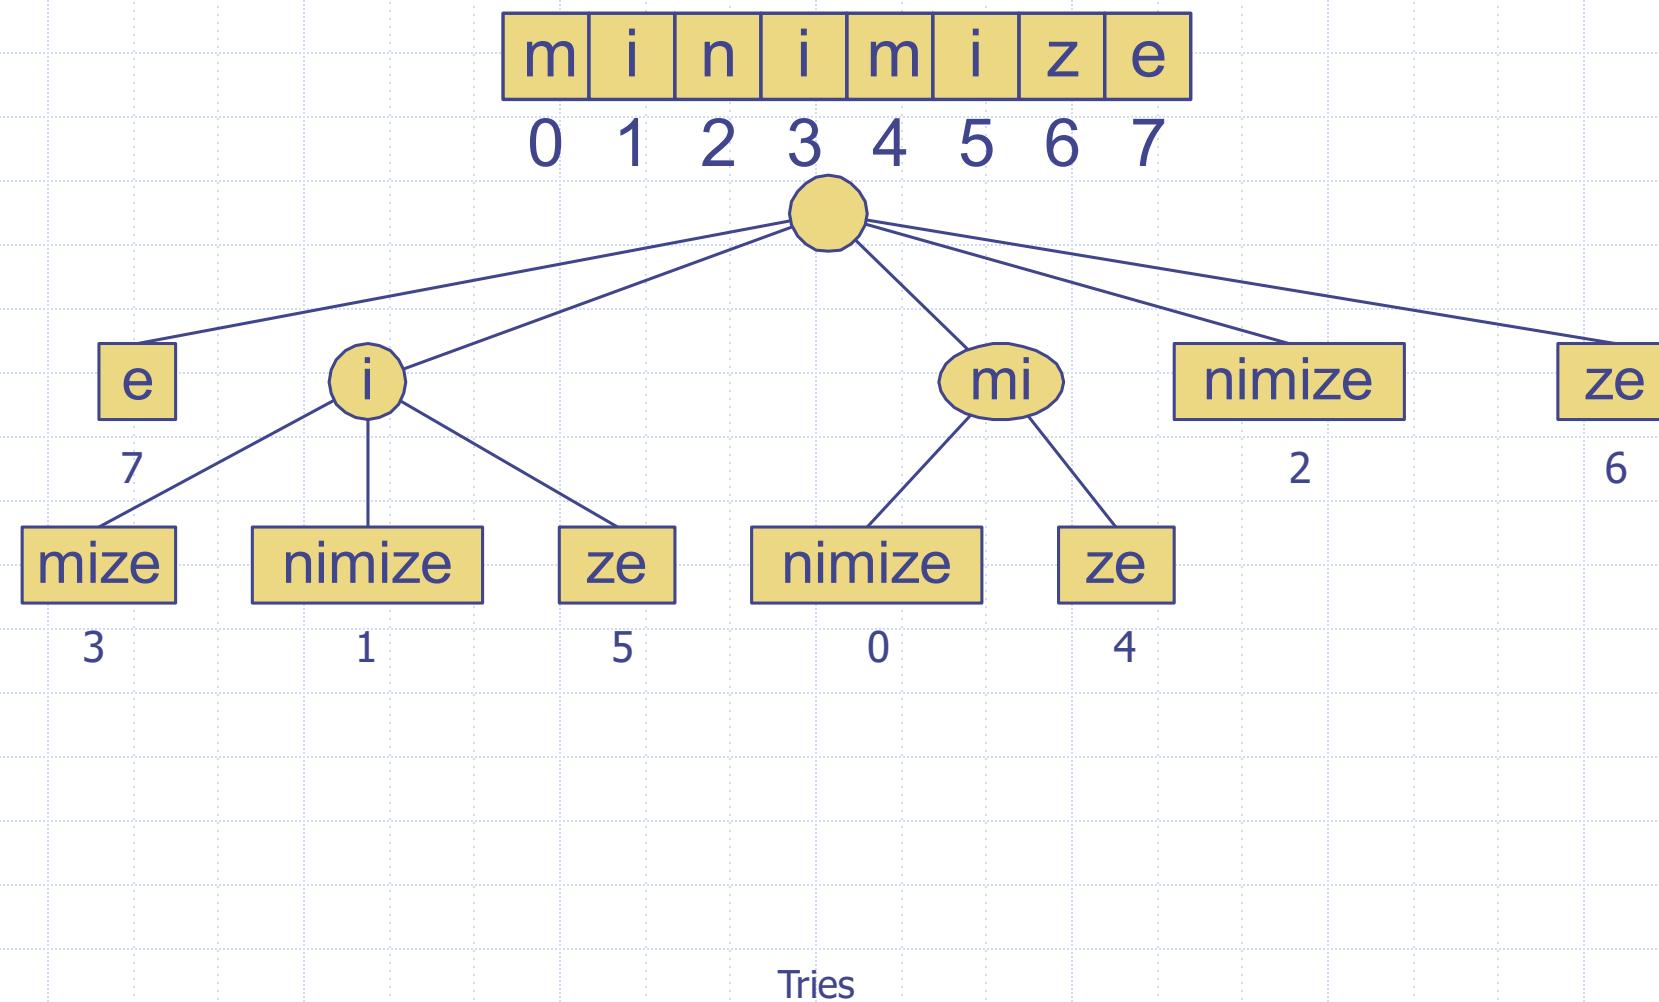
Example



Example (2)



Example (3)



Constructing a Suffix Trie

- $S[1..n]$ is the string
- start with a single edge for S
- enter the edges for the suffix $S[i..n]$ where i goes from 2 to n
 - Starting at the root node find the longest part from the root whose label matches a prefix of $S[i..n]$. At some point, no further matches are possible
 - ◆ If the point is at a node, then denote this node by w
 - ◆ If it is in the middle of an edge, then insert a new node called w , at this point
 - ◆ create a new edge running from w to a new leaf labeled $S[i..n]$

Complexity- $O(n^2)$

Dynamic Programming



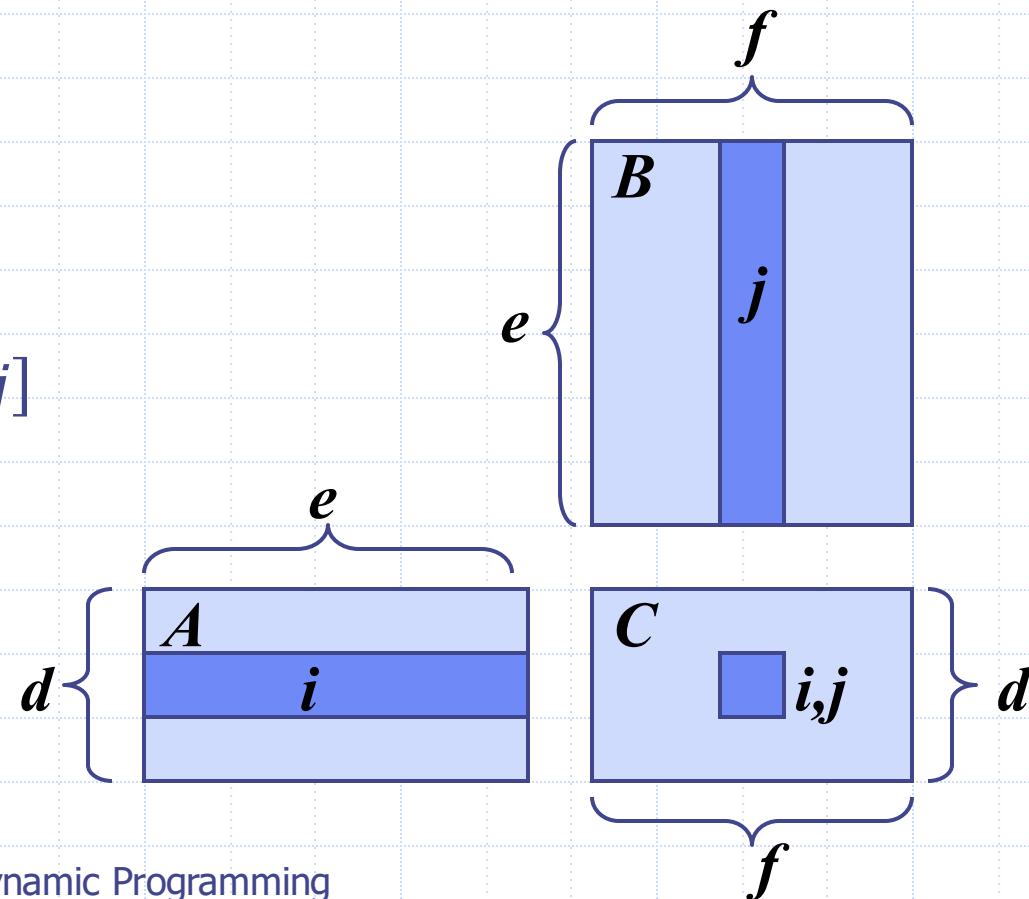
Motivating Example

- Fibonacci number computation
 - Recursive definition
 - ◆ $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$
 - Complexity – $O(\Phi^n)$

Matrix Chain-Products

- Dynamic Programming is a general algorithm design paradigm.
 - Matrix Chain-Products
- Review: Matrix Multiplication.
 - $C = A * B$
 - A is $d \times e$ and B is $e \times f$
 - $O(def)$ time

$$C[i, j] = \sum_{k=0}^e A[i, k]B[k, j]$$



Matrix Chain-Products

□ Matrix Chain-Product:

- Compute $A = A_0 * A_1 * \dots * A_{n-1}$
- A_i is $d_i \times d_{i+1}$
- Problem: How to parenthesize?

□ Example

- B is 3×100
- C is 100×5
- D is 5×5
- $(B*C)*D$ takes $1500 + 75 = 1575$ ops
- $B*(C*D)$ takes $1500 + 2500 = 4000$ ops

An Enumeration Approach

- **Matrix Chain-Product Alg.:**

- Try all possible ways to parenthesize $A = A_0 * A_1 * \dots * A_{n-1}$
- Calculate number of ops for each one
- Pick the one that is best

- Running time:

- This is **exponential!**
- It is called the Catalan number, and it is almost 4^n .
- This is a terrible algorithm!

A Greedy Approach

- Idea #1: repeatedly select the product that uses (up) the most operations.
- Counter-example:
 - A is 10×5
 - B is 5×10
 - C is 10×5
 - D is 5×10
 - Greedy idea #1 gives $(A*B)*(C*D)$, which takes $500+1000+500 = 2000$ ops
 - $A*((B*C)*D)$ takes $500+250+250 = 1000$ ops

Another Greedy Approach

- Idea #2: repeatedly select the product that uses the fewest operations.
- Counter-example:
 - A is 101×11
 - B is 11×9
 - C is 9×100
 - D is 100×99
 - Greedy idea #2 gives $A * ((B * C) * D)$, which takes $109989 + 9900 + 108900 = 228789$ ops
 - $(A * B) * (C * D)$ takes $9999 + 89991 + 89100 = 189090$ ops
- The greedy approach is not giving us the optimal value.

A “Recursive” Approach



- Define **subproblems**:
 - Find the best parenthesization of $A_i * A_{i+1} * \dots * A_j$.
 - Let $N_{i,j}$ denote the number of operations done by this subproblem.
 - The optimal solution for the whole problem is $N_{0,n-1}$.
- **Subproblem optimality**: The optimal solution can be defined in terms of optimal subproblems
 - There has to be a final multiplication (root of the expression tree) for the optimal solution.
 - Say, the final multiply is at index i : $(A_0 * \dots * A_i) * (A_{i+1} * \dots * A_{n-1})$.
 - Then the optimal solution $N_{0,n-1}$ is the sum of two optimal subproblems, $N_{0,i}$ and $N_{i+1,n-1}$ plus the time for the last multiply.
 - If the global optimum did not have these optimal subproblems, we could define an even better “optimal” solution.

A Characterizing Equation

- The global optimal has to be defined in terms of optimal subproblems, depending on where the final multiply is at.
- Let us consider all possible places for that final multiply:
 - Recall that A_i is a $d_i \times d_{i+1}$ dimensional matrix.
 - So, a characterizing equation for $N_{i,j}$ is the following:

$$N_{i,j} = \min_{i \leq k < j} \{N_{i,k} + N_{k+1,j} + d_i d_{k+1} d_{j+1}\}$$

- Note that subproblems are not independent--the **subproblems overlap**.

A Dynamic Programming Algorithm

- Since subproblems overlap, we don't use recursion.
- Instead, we construct optimal subproblems "bottom-up."
- $N_{i,i}$'s are easy, so start with them
- Then do length 2,3,... subproblems, and so on.
- The running time is $O(n^3)$

Algorithm *matrixChain(S)*:

Input: sequence S of n matrices to be multiplied

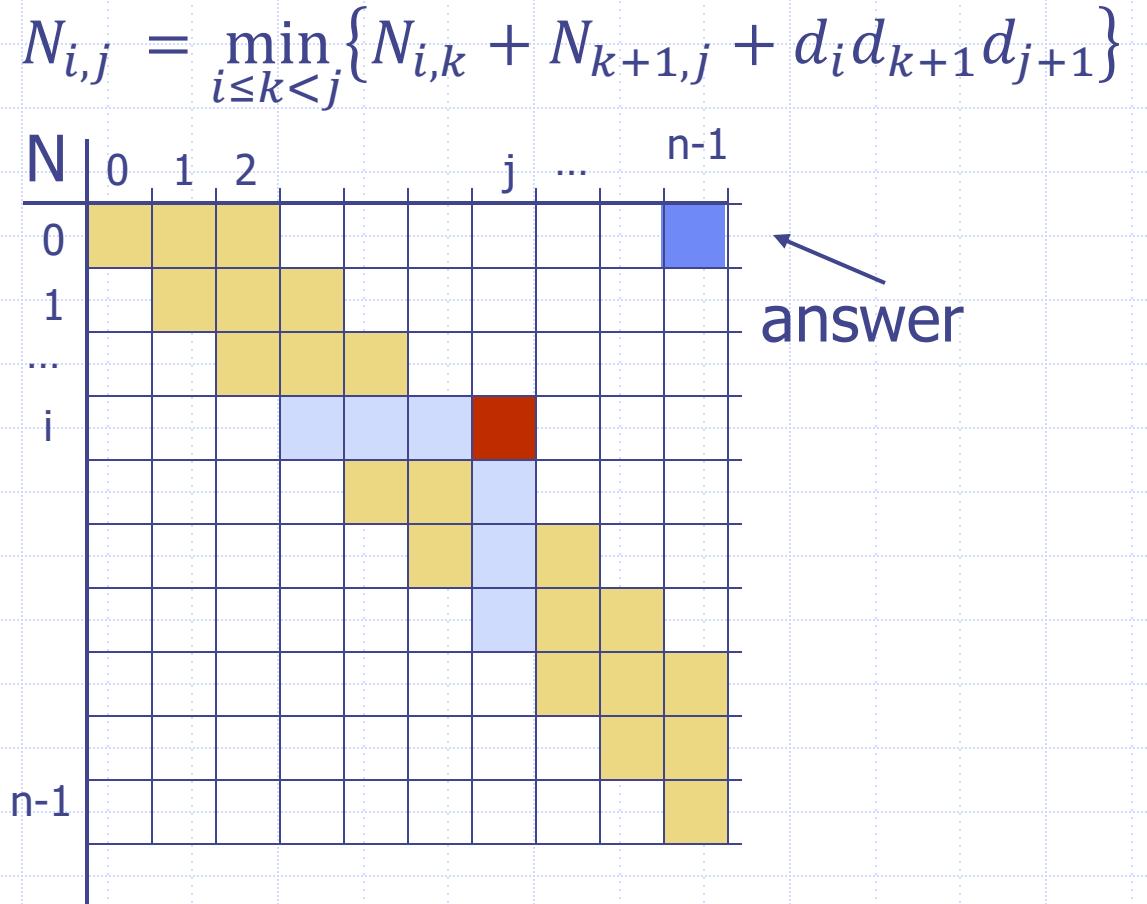
Output: number of operations in an optimal parenthization of S

```
for  $i \leftarrow 1$  to  $n-1$  do  
   $N_{i,i} \leftarrow 0$   
for  $b \leftarrow 1$  to  $n-1$  do  
  for  $i \leftarrow 0$  to  $n-b-1$  do  
     $j \leftarrow i+b$   
     $N_{i,j} \leftarrow +\infty$   
    for  $k \leftarrow i$  to  $j-1$  do  
       $N_{i,j} \leftarrow \min\{N_{i,j}, N_{i,k} + N_{k+1,j}$   
       $+ d_i d_{k+1} d_{j+1}\}$ 
```

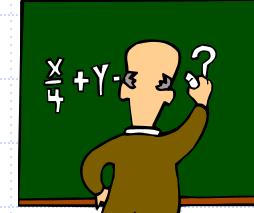
A Dynamic Programming Algorithm Visualization



- The bottom-up construction fills in the N array by diagonals
- $N_{i,j}$ gets values from previous entries in i-th row and j-th column
- Filling in each entry in the N table takes $O(n)$ time.
- Total run time: $O(n^3)$
- Getting actual parenthesization can be done by remembering “k” for each N entry



The General Dynamic Programming Technique



- Applies to a problem that at first seems to require a lot of time (possibly exponential), provided we have:
 - **Simple subproblems:** the subproblems can be defined in terms of a few variables, such as j , k , l , m , and so on.
 - **Subproblem optimality:** the global optimum value can be defined in terms of optimal subproblems
 - **Subproblem overlap:** the subproblems are not independent, but instead they overlap (hence, should be constructed bottom-up).

Subsequences

- A ***subsequence*** of a character string $x_0x_1x_2\dots x_{n-1}$ is a string of the form $x_{i_1}x_{i_2}\dots x_{i_k}$, where $i_j < i_{j+1}$.
- Not the same as substring!
- Example String: ABCDEFGHIJK
 - Subsequence: ACEGIJK
 - Subsequence: DFGHK
 - Not subsequence: DAGH

The Longest Common Subsequence (LCS)

Problem

- Given two strings X and Y, the longest common subsequence (LCS) problem is to find a longest subsequence common to both X and Y
- Has applications to DNA similarity testing (alphabet is {A,C,G,T})
- Example: ABCDEFG and XZACKDFWGH have ACDFG as a longest common subsequence

A Poor Approach to the LCS Problem

- A Brute-force solution:
 - Enumerate all subsequences of X
 - Test which ones are also subsequences of Y
 - Pick the longest one.
- Analysis:
 - If X is of length n, then it has 2^n subsequences
 - This is an exponential-time algorithm!