


Minimum Spanning Trees



Minimum Spanning Trees

Spanning subgraph

- Subgraph of a graph G containing all the vertices of G

Spanning tree


- Spanning subgraph that is itself a (cycle free) tree

Minimum spanning tree (MST)

- Spanning tree of a weighted graph with minimum total edge weight

Applications

- Communications networks
- Transportation networks




Cycle Property

Cycle Property:


- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $\text{weight}(f) \leq \text{weight}(e)$

Proof:

- By contradiction
- If $\text{weight}(f) > \text{weight}(e)$ we can get a spanning tree of smaller weight by replacing e with f




Replacing f with e yields a better spanning tree



Cuts

- A cut in G is a partition of the vertex set into two parts
 - Example:
 - ◆ $U = \{a, e, b\}$
 - ◆ $V = \{c, d, f\}$
 - Number of possible cuts
 - ◆ $2^{n-1} - 1$
- Edges in the cut
 - edges that connect vertices of one partition to vertices of the other partition.




Partition Property

Partition Property:


- Consider a partition of the vertices of G into subsets U and V
- Let e be an edge of minimum weight across the partition
- There is a minimum spanning tree of G containing edge e

Proof:

- Let T be an MST of G
- If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition
- By the cycle property,
 $weight(f) \leq weight(e)$
- Thus, $weight(f) = weight(e)$
- We obtain another MST by replacing f with e




Replacing f with e yields another MST



Prim-Jarnik's Principle


- ❑ Similar to Dijkstra's algorithm
- ❑ Utilizes the Partition Property
- ❑ We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s




Prim-Jarnik's Algorithm - Implementation

- Similar to Dijkstra's algorithm
- We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s
- We store with each vertex v label $d(v)$ representing the smallest weight of an edge connecting v to a vertex in the cloud
- At each step:
 - We add to the cloud the vertex u outside the cloud with the smallest distance label
 - We update the labels of the vertices adjacent to u

Example



Example (contd.)



Prim-Jarnik Pseudo-code

Algorithm PrimJarnik(G):

Input: An undirected, weighted, connected graph G with n vertices and m edges

Output: A minimum spanning tree T for G

Pick any vertex s of G

$D[s] = 0$

for each vertex $v \neq s$ **do**

$D[v] = \infty$

Initialize $T = \emptyset$.

Initialize a priority queue Q with an entry $(D[v], (v, \text{None}))$ for each vertex v , where $D[v]$ is the key in the priority queue, and (v, None) is the associated value.

while Q is not empty **do**

$(u, e) = \text{value returned by } Q.\text{remove_min}()$

 Connect vertex u to T using edge e .

for each edge $e' = (u, v)$ such that v is in Q **do**

 {check if edge (u, v) better connects v to T }

if $w(u, v) < D[v]$ **then**

$D[v] = w(u, v)$

 Change the key of vertex v in Q to $D[v]$.

 Change the value of vertex v in Q to (v, e') .

return the tree T


Analysis

- Graph operations
 - We cycle through the incident edges once for each vertex
- Label operations
 - We set/get the distance, parent and locator labels of vertex z $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex w in the priority queue is modified at most $\deg(w)$ times, where each key change takes $O(\log n)$ time
- Prim-Jarnik's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time is $O(m \log n)$ since the graph is connected

Kruskal's Principle

□ Greedy Algorithm


- greedily add the edge with minimum weight to the MST
- ensure that the added edge does not form a cycle



Kruskal's Algorithm


- Maintain a partition of the vertices into clusters
 - Initially, single-vertex clusters
 - Keep an MST for each cluster
 - Merge “closest” clusters and their MSTs
- A priority queue stores the edges outside clusters
 - Key: weight
 - Element: edge
- At the end of the algorithm
 - One cluster and one MST

Example



Minimum Spanning Trees

Example (contd.)



Algorithm Kruskal(G):

Input: A simple connected weighted graph G with n vertices and m edges

Output: A minimum spanning tree T for G

for each vertex v in G do

Define an elementary cluster $C(v) = \{v\}$.

Initialize a priority queue Q to contain all edges in G , using the weights as keys.

$T = \emptyset$ { T will ultimately contain the edges of the MST}

while T has fewer than $n - 1$ edges **do**

(u, v) = value returned by $Q.remove_min()$

Let $C(u)$ be the cluster containing u , and let $C(v)$ be the cluster containing v .

if $C(u) \neq C(v)$ then

Add edge (u, v) to T .

Merge $C(u)$ and $C(v)$ into one cluster.

return tree T

Correctness of Kruskal's Algorithm

- Let all the edges in the graph have distinct weights
- Let the tree obtained from Kruskal's algorithm contain the edges $g_1 < g_2 < g_3 \dots < g_{n-1}$
- Let the Optimum MST have the edges $f_1 < f_2 < f_3 \dots < f_{n-1}$
 - Let i be the first edge that is different between g 's and f 's.
 - ♦ case1: $g_i < f_i$
 - ♦ case2: $f_i < g_i$

Is the MST unique?

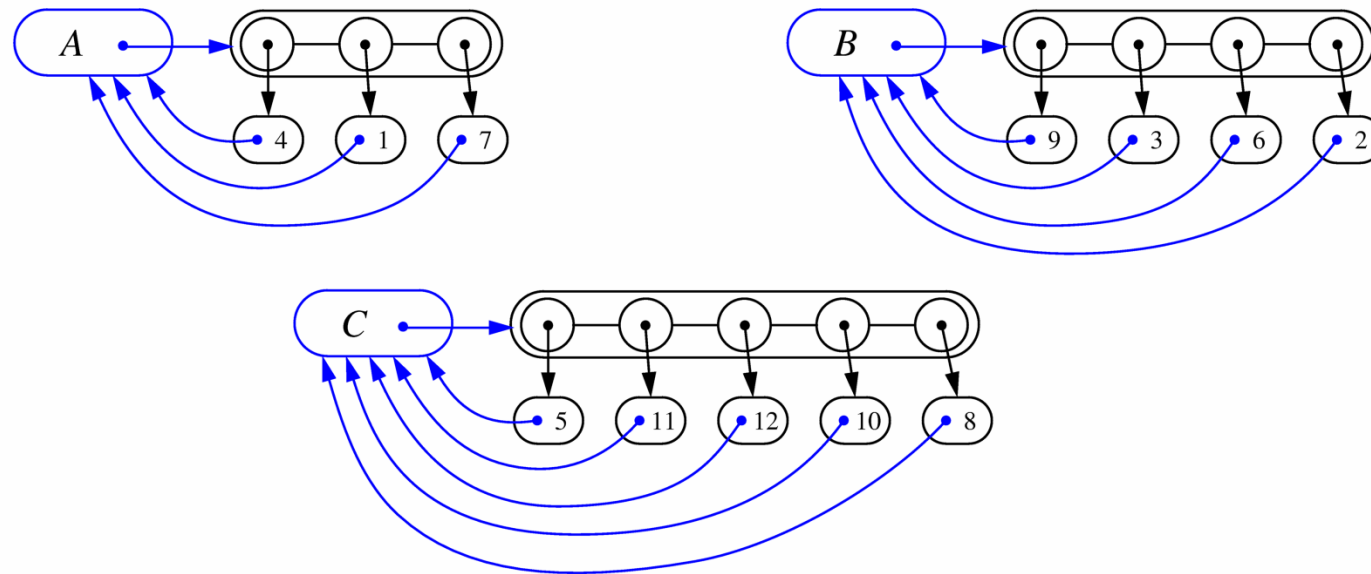
- yes, if the edge weights are distinct.
- no, otherwise.

Data Structure for Kruskal's Algorithm

- The algorithm maintains a forest of trees
- A priority queue extracts the edges by increasing weight
- An edge is accepted if it connects distinct trees
- We need a data structure that maintains a **partition**, i.e., a collection of disjoint sets, with operations:
 - **makeSet**(u): create a set consisting of u
 - **find**(u): return the set storing u
 - **union**(A, B): replace sets A and B with their union
- **Total time**
 - $O(m \log n + \text{time for Union} * n + \text{time for Find} * m)$

List-based Implementation

- ❑ Each set is stored in a sequence represented with a linked-list
- ❑ Each node should store an object containing the element and a reference to the set name



Analysis of List-based Representation

- ◆ Union

- ◆ merging of two linked lists
- ◆ $O(n)$

- ◆ Find


- ◆ search through the list
- ◆ $O(n)$

- ◆ Total time

- ◆ $O(m \log n) + O(n^2) + O(mn)$.



Tree-based Implementation

- ❑ Each element is stored in a node, which contains a pointer to a **set** name
- ❑ A node v whose set pointer points back to v is also a set name
- ❑ Each set is a tree, rooted at a node with a self-referencing set pointer
- ❑ For example: The sets “1”, “2”, and “5”:



Union-Find Operations

- ❑ To do a **union**, simply make the root of one tree point to the root of the other
- ❑ To do a **find**, follow set-name pointers from the starting node until reaching a node whose set-name pointer refers back to itself
- ❑ Amortized time is $O(n \log^* n)$





The End!

Course Schedule

week 1 - Introduction OOP, Arrays, and Linked Lists
week 2 - Analysis Tools and Recursion
week 3 - Stacks and Queues
week 4 - Tree Structures and Heaps
week 5 - Sorting and Priority Queues
week 6 - Maps, Hash Tables
week 7 - Search Tree Structures I, **Quiz 1**
week 8 - Search Tree Structures II
week 9 - Search Tree Structures III
week 10 - Buffer week
week 11 - String Processing
week 12 - Dynamic Programming, **Quiz 2**
week 13 - Graphs I
week 14 - Graphs II
week 15 - Graphs III
week 16 - Buffer week

Arrays and Linked Lists

- Chapter 3, Sections 3.1-3.4
- Arrays
 - Insertion sort
- Singly Linked Lists
- Doubly Linked Lists

Analysis of Algorithms

- ❑ Chapter 4, Sections 4.1-4.4
- ❑ Writing and analyzing pseudo codes
- ❑ Big Oh Notation
 - asymptotic algorithm analysis
- ❑ Big Omega and Big Theta notations
- ❑ Proof methods
- ❑ mathematical functions and their relationships

Recursion

- ❑ Chapter 5, Section 5.1-5.5
- ❑ Defining a recursion pattern
- ❑ Analysis of recursion
- ❑ Types of recursion
 - Linear, binary and multiple
- ❑ Pitfalls of recursion

Stacks and Queues

- ❑ Chapter 6, Sections 6.1-6.3
- ❑ ADT
- ❑ Stacks
 - ADT definition
 - Implementations
 - Applications
- ❑ Queues
 - ADT definition
 - Implementations
 - Applications
 - Circular Queues

Tree Structures

- Chapter 8, Sections 8.1-8.4
- Generic Tree ADT
- Binary Trees
 - properties
- Tree Representations
- Tree Traversals
 - preorder
 - postorder
 - inorder
 - breadth first

Heaps and Priority Queues

- Chapter 9, Sections 9.1-9.4
- Priority Queue
 - ADT Definition
 - Implementation using sorted/unsorted lists
- Heaps
 - Definition
 - Priority Queue as a heap
 - Analysis of Upheap and Downheap procedures
 - Heap construction methods
 - ◆ analysis
- Sorting with a priority queue
 - selection sort
 - insertion sort
 - heap sort

Searching and Search Tree Structures (chapter 11)

- Binary Search Trees
 - insertion and deletion operations
- Balanced Search Trees
- AVL Trees
 - operations
- (2,4) Trees
 - multi-way search trees

Hash Tables and Maps

- Chapter 10, Sections 10.1-10.3
- Map
 - ADT
 - applications
- Hashing
 - Hash functions
 - Collision handling schemes
 - load factors and efficiency

String Processing and Dynamic Programming (Chapter 12)

- Pattern Matching Algorithms
 - The Boyer-Moore Algorithm
 - ◆ Bad character and good suffix rules
- Tries
 - Standard Tries
 - Compressed Tries
 - Suffix Tries
 - ◆ Prefix Matching
- Dynamic Programming
 - Methodology

Sorting (Chapter 13)

- ❑ Divide and Conquer Paradigm
- ❑ Merge Sort
 - algorithm
 - complexity analysis
- ❑ Quick Sort/Randomized Quick Sort
 - algorithm
 - complexity analysis
 - In-Place partitioning
- ❑ Radix Sort
 - Radix Exchange Sort
 - Straight Radix Sort
 - ◆ Stable way
- ❑ Bucket sort

Graphs (Chapter 14)

- Graphs
 - Terminology
- Data Structures for representing graphs
 - edge list, adjacency list, adjacency matrix
- Graph Traversals
 - BFS, DFS
- Directed Graphs
- Transitive Closure
- Directed Acyclic Graphs
 - topological order
- Shortest Paths
 - weighted graphs, Dijkstra's algorithm, Bellman-ford algorithm
- Minimum Spanning Trees
 - Prim-Jarnik Algorithm, Kruskal's Algorithm, Data structure for Kruskal's algorithm.

Good Luck



**KEEP
CALM
MAY THE
FORCE BE
WITH YOU**